

## Spark en EMR y Google Colab

```
(1) [1]: spark
... SparkSession - in-memory
SparkContext
SparkUI
Version
  v4.0.1
Master
  Local[*]
AppName
  pyspark-shell

(2) [1]: #configuración en google colab de spark y pyspark
from google.colab import drive
drive.mount('/content/gdrive')
Mounted at /content/gdrive

(3) [2]: !instalar java y spark
apt-get install openjdk-17-jdk-headless -qq > /dev/null
wget -q https://downloads.apache.org/spark/spark-4.0.1/spark-4.0.1-bin-hadoop3.tgz
tar xf spark-4.0.1-bin-hadoop3.tgz
!pip install -q findspark

(4) [3]: import os
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-17-openjdk-amd64"
os.environ["SPARK_HOME"] = "/content/spark-4.0.1-bin-hadoop3"

(5) [4]: import findspark
findspark.init()
from pyspark.sql import SparkSession
spark = SparkSession.builder.master("local[*]").getOrCreate()
sc = spark.sparkContext

(6) [5]: # Load csv dataset
df=spark.read.csv('sample_data/california_housing_test.csv',inferSchema=True,header=True)
# desde S3
# dfspark.read.csv('s3a://bucle_name/datasets/sample_data.csv',inferSchema=True,header=True)

(7) [6]: #columns of dataframe
df

(8) [7]: #dataframe[longitude: double, latitude: double, housing_median_age: double, total_rooms: double, total_bedrooms: double, population: double, households: double, median_income: double, median_house_value: double]

(9) [8]: #check number of columns
len(df.columns)

(10) [9]: 9

(11) [10]: #number of records in dataframe
df.count()
```

```
(1) [1]: spark
... SparkSession - in-memory
SparkContext
SparkUI
Version
  v4.0.1
Master
  Local[*]
AppName
  pyspark-shell

(2) [2]: !instalar java y spark
apt-get install openjdk-17-jdk-headless -qq > /dev/null
wget -q https://downloads.apache.org/spark/spark-4.0.1/spark-4.0.1-bin-hadoop3.tgz
tar xf spark-4.0.1-bin-hadoop3.tgz
!pip install -q findspark

(3) [3]: import os
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-17-openjdk-amd64"
os.environ["SPARK_HOME"] = "/content/spark-4.0.1-bin-hadoop3"

(4) [4]: import findspark
findspark.init()
from pyspark.sql import SparkSession
spark = SparkSession.builder.master("local[*]").getOrCreate()
sc = spark.sparkContext

(5) [5]: # Load csv Dataset
df=spark.read.csv('sample_data/california_housing_test.csv',inferSchema=True,header=True)
# desde S3
# dfspark.read.csv('s3a://bucle_name/datasets/sample_data.csv',inferSchema=True,header=True)
```

google-colab-setup-PySpark.ipynb

Archivo Editar Ver Insertar Entorno de ejecución Herramientas Ayuda

mandos + Código + Texto ▶ Ejecutar todas

Índice

Data Processing using Pyspark

+ Sección

```
[1]: 'longitude',
      'housing_median_age',
      'total_rooms',
      'total_bedrooms',
      'population',
      'households',
      'median_income',
      'median_house_value']

[2]: #shape of dataset
print((df.count()),len(df.columns))

(3800, 9)

[3]: #printSchema
df.printSchema()

root
 |-- longitude: double (nullable = true)
 |-- latitude: double (nullable = true)
 |-- housing_median_age: double (nullable = true)
 |-- total_rooms: double (nullable = true)
 |-- total_bedrooms: double (nullable = true)
 |-- population: double (nullable = true)
 |-- households: double (nullable = true)
 |-- median_income: double (nullable = true)
 |-- median_house_value: double (nullable = true)

[4]: #first few rows of dataframe
df.show(5)

+-----+-----+-----+-----+-----+-----+-----+-----+
|longitude|latitude|housing_median_age|total_rooms|total_bedrooms|population|households|median_income|median_house_value|
+-----+-----+-----+-----+-----+-----+-----+-----+
| -122.05 | 37.37 | 27.0 | 3885.0 | 661.0 | 1537.0 | 606.0 | 6.6085 | 344780.0 |
| -118.3 | 34.26 | 43.0 | 1510.0 | 310.0 | 899.0 | 277.0 | 3.599 | 176500.0 |
| -117.81 | 33.78 | 27.0 | 3589.0 | 580.0 | 1484.0 | 495.0 | 5.7934 | 279500.0 |
| -117.36 | 33.87 | 28.0 | 3707.0 | 15.0 | 487.0 | 132.0 | 5.3599 | 330000.0 |
| -119.67 | 36.33 | 19.0 | 1241.0 | 244.0 | 856.0 | 237.0 | 2.9375 | 81700.0 |
+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows
```

google-colab-setup-PySpark.ipynb

Archivo Editar Ver Insertar Entorno de ejecución Herramientas Ayuda

mandos + Código + Texto ▶ Ejecutar todas

Índice

Data Processing using Pyspark

+ Sección

```
[1]: SparkContext
Spark_UU
Version
v4.0.1
Master
local[*]
AppName
pyspark-shell

[2]: spark
SparkSession - in-memory
SparkContext
Spark_UU
Version
v4.0.1
Master
local[*]
AppName
pyspark-shell

[3]: #configuración en google colab de spark y pyspark
from google.colab import drive
drive.mount('/content/gdrive')

[4]: #instalar java y spark
!apt-get install openjdk-17-jdk-headless -qq > /dev/null
!wget -q https://downloads.apache.org/spark/spark-4.0.1/spark-4.0.1-bin-hadoop3.tgz
!tar xf spark-4.0.1-bin-hadoop3.tgz
!pip install -q findspark
```

google-colab-setup-PySpark2AWS-S3.ipynb

```
[1]: #configuración en google colab de spark y pyspark
from google.colab import drive
drive.mount('/content/gdrive')
Mounted at /content/gdrive

[2]: #Instalar java y spark
!apt-get install openjdk-17-jdk-headless -qq > /dev/null
!wget -q https://downloads.apache.org/spark/spark-3.5.7/spark-3.5.7-bin-hadoop3.tgz
!tar xf spark-3.5.7-bin-hadoop3.tgz
!pip install -q findspark

[3]: import os
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-17-openjdk-amd64"
os.environ["SPARK_HOME"] = "/content/spark-3.5.7-bin-hadoop3"

[4]: import findspark
findspark.init()

[5]: !ls -p /content/jars
!wget -q https://repo.maven.org/maven2/org/apache/hadoop/hadoop-mr/3.3.4/hadoop-mr-3.3.4.jar -P /content/jars
!wget -q https://repo.maven.org/maven2/com/amazonaws/aws-java-sdk-bundle/1.12.367/aws-java-sdk-bundle-1.12.367.jar -P /content/jars

[6]: from pyspark.sql import SparkSession
jars = "/content/jars/hadoop-aws-3.3.4.jar,/content/jars/aws-java-sdk-bundle-1.12.367.jar"

spark = SparkSession.builder \
    .appName("S3connection") \
    .master("local[*]") \
    .config("spark.jars", jars) \
    .config("fs.s3a.access.key", "SIAISN2J3N2F266MEUHF") \
    .config("fs.s3a.secret.key", "IG7BVWKL1J7ZnH0Bxtgyspf2dvHkdw6N6gkIE9") \
    .config("fs.s3a.session.token", "IQoJb3ppZ2luXZV2ENf//////////EaCXzLXd13QmL1JHMCEUCIBzqr642+hxzFbdB6lPyvNt3zAuFIqoiJOjDvY2MPBAIEuRwOnpshrm3lr/1xALenu3wTAHgwD2JdxmInKbUX/sqr") \
    .config("spark.hadoop.fs.s3a.impl", "org.apache.hadoop.fs.s3a.S3AFileSystem") \
    .config("spark.hadoop.fs.s3a.path.style.access", "true") \
    .config("spark.hadoop.fs.s3a.endpoint", "ss.amazonaws.com") \
```

google-colab-setup-PySpark2AWS-S3.ipynb

```
[1]: df=spark.read.csv('s3a://emrsmuel/datasets/sample_data.csv',inferSchema=True,header=True)
#columns of dataframe
df

[2]: DataFrame(ratings: int, age: int, experience: double, family: int, mobile: string)

[3]: #check number of columns
len(df.columns)
5

[4]: #number of records in dataframe
df.count()
5

[5]: #shape of dataset
print(df.count(),len(df.columns))
(35, 5)

[6]: df.printSchema()
df.printSchema()

*** root
|-- ratings: integer (nullable = true)
|-- age: integer (nullable = true)
|-- experiences double (nullable = true)
|-- family: integer (nullable = true)
|-- mobile: string (nullable = true)

[7]: #first few rows of dataframe
df.show(5)

+---+---+---+---+---+
|ratings|age|experience|family|mobile|
+---+---+---+---+---+
| 3| 32| 9.0| 3| Vivo|
| 3| 32| 13.0| 3| Apple|
| 4| 22| 2.5| 0|Samsung|
| 4| 37| 16.5| 4| Apple|
| 5| 27| 9.0| 1| MI|
+---+---+---+---+---+
```

Data\_processing\_using\_PySpark.ipynb

Archivo Editar Ver Insertar Entorno de ejecución Herramientas Ayuda

Comandos + Código + Texto | ▶ Ejecutar todas ▾

```
[✓ 0s] #df=spark.read.csv('s3a://emrsamuel/datasets/sample_data.csv',inferSchema=True,header=True)
[✓ 0s] [8] #columns of dataframe
[✓ 0s] df.columns
[✓ 0s] ['ratings', 'age', 'experience', 'family', 'mobile']
[✓ 0s] [9] #check number of columns
[✓ 0s] len(df.columns)
[✓ 0s] 5
[✓ 0s] [10] #number of records in dataframe
[✓ 0s] df.count()
[✓ 0s] 33
[✓ 0s] [11] #shape of dataset
[✓ 0s] print((df.count(),len(df.columns)))
[✓ 0s] (33, 5)
[✓ 0s] [12] #printSchema
[✓ 0s] df.printSchema()
[✓ 0s] root
[✓ 0s] |-- ratings: integer (nullable = true)
[✓ 0s] |-- age: integer (nullable = true)
[✓ 0s] |-- experience: double (nullable = true)
[✓ 0s] |-- family: integer (nullable = true)
[✓ 0s] |-- mobile: string (nullable = true)
[✓ 0s] [13] #first few rows of dataframe
[✓ 0s] df.show(5)
[✓ 0s]
+---+---+---+---+
|ratings|age|experience|family| mobile|
+---+---+---+---+
| 3| 32| 9.0| 3| Vivo|
| 3| 27| 13.0| 3| Apple|
| 4| 22| 2.5| 0|Samsung|
| 4| 37| 16.5| 4| Apple|
| 5| 27| 9.0| 1| MI|
+---+---+---+---+
only showing top 5 rows
```

Variables Terminal

DATA PROCESSING USING PYSPARK

Archivo Editar Ver Insertar Entorno de ejecución Herramientas Ayuda

Comandos + Código + Texto | Ejecutar todas

[14] ✓ 0s #select only 2 columns  
df.select('age','mobile').show(5)

age	mobile
32	Vivo
27	Apple
22	Samsung
37	Apple
27	MI

only showing top 5 rows

[15] ✓ 2s #info about dataframe  
df.describe().show()

summary	ratings	age	experience	family	mobile
count	33	33	33	33	33
mean	3.57575757575757	30.4848484848484	10.3030303030303	1.81818181818181	NULL
stddev	1.1188806636071336	6.18527087180309	6.770731351213326	1.8448330794164254	NULL
min	1	22	2.5	0	Apple
max	5	42	23.0	5	Vivo

[16] ✓ 0s from pyspark.sql.types import StringType,DoubleType,IntegerType

[17] ✓ 0s #add column  
df.withColumn("age\_after\_10\_yrs", (df["age"]+10)).show(10,False)

ratings	age	experience	family	mobile	age_after_10_yrs
3	32	9.0	3	Vivo	42
3	27	13.0	3	Apple	37
4	22	2.5	0	Samsung	32
4	37	16.5	4	Apple	47
5	27	9.0	1	MI	37
4	27	9.0	0	Oppo	37
5	37	23.0	5	Vivo	47
5	37	23.0	5	Samsung	47
3	22	2.5	0	Apple	32
3	27	6.0	0	MI	37

only showing top 10 rows

Data\_processing\_using\_PySpark.ipynb

Archivo Editar Ver Insertar Entorno de ejecución Herramientas Ayuda

Comandos + Código + Texto ▶ Ejecutar todas

[21] ✓ 0 s

```
#filter the records
df.filter(df['mobile']=='Vivo').select('age','ratings','mobile').show()
```

	5  37	13.0	1  Vivo
...	4  37	6.0	0  Vivo

[22] ✓ 0 s

```
#filter the multiple conditions
df.filter(df['mobile']=='Vivo').filter(df['experience'] >10).show()
```

	ratings age experience family mobile
	5  37  23.0  5  Vivo
	5  37  13.0  1  Vivo

[23] ✓ 0 s

```
#filter the multiple conditions
df.filter((df['mobile']=='Vivo')&(df['experience'] >10)).show()
```

	ratings age experience family mobile
	5  37  23.0  5  Vivo
	5  37  13.0  1  Vivo

[24] ✓ 1 s

```
#Distinct Values in a column
df.select('mobile').distinct().show()
```

	mobile
	MI

Variables Terminal

⚠ Data\_processing\_using\_PySpark.ipynb ⚡ 🔍

Archivo Editar Ver Insertar Entorno de ejecución Herramientas Ayuda

Comandos + Código + Texto ▶ Ejecutar todas

```
[25] ✓ 0 s
#distinct value count
df.select('mobile').distinct().count()

5

[26] ✓ 0 s
▶ df.groupBy('mobile').count().show(5,False)

+---+---+
|mobile |count|
+---+---+
|MI    |8   |
|Oppo  |7   |
|Samsung|6   |
|Vivo   |5   |
|Apple  |7   |
+---+---+


[27] ✓ 0 s
# Value counts
df.groupBy('mobile').count().orderBy('count',ascending=False).show(5,False)

+---+---+
|mobile |count|
+---+---+
|MI      |8   |
|Oppo    |7   |
|Apple   |7   |
|Samsung|6   |
|Vivo    |5   |
+---+---+


[28] ✓ 0 s
# Value counts
df.groupBy('mobile').mean().show(5,False)

+-----+-----+-----+-----+-----+
|mobile |avg(ratings) |avg(age)    |avg(experience)|avg(family)  |
+-----+-----+-----+-----+-----+
|MI     |3.5        |30.125     |10.1875      |1.375       |
|Oppo   |2.857142857142857|28.428571428571427|10.357142857142858|1.4285714285714286|
|Samsung|4.1666666666666667|28.666666666666668|8.666666666666666|1.8333333333333333|
|Vivo   |4.2        |36.0       |11.4         |1.8         |
|Apple  |3.4285714285714284|30.571428571428573|11.0         |2.7142857142857144|
+-----+-----+-----+-----+-----+
```

Variables Terminal

⚠ Data\_processing\_using\_PySpark.ipynb ⚡ ↗

Archivo Editar Ver Insertar Entorno de ejecución Herramientas Ayuda

mandos + Código ▾ + Texto | ▶ Ejecutar todas ▾

```
[43] [✓ 0s] #duplicate values
df.count()

33

[44] [✓ 0s] #drop duplicate values
df=df.dropDuplicates()

[45] [✓ 0s] #validate new count
df.count()

26

[46] [✓ 0s] #drop column of dataframe
df_new=df.drop('mobile')

[47] [✓ 0s] df_new.show(10)

+---+---+---+---+
|ratings|age|experience|family|
+---+---+---+---+
|      3| 32|      9.0|    3|
|      4| 22|      2.5|    0|
|      5| 27|      6.0|    0|
|      4| 22|      6.0|    1|
|      3| 27|      6.0|    0|
|      2| 32|     16.5|    2|
|      4| 27|      9.0|    0|
|      2| 27|      9.0|    2|
|      3| 37|     16.5|    5|
|      4| 27|      6.0|    1|
+---+---+---+---+
only showing top 10 rows

[48] [✓ 0s] # saving file to csv
```

ID	YARN Application ID	Kind	State	Spark UI	Driver log	User	Current session?
3	application_1763344109923_0006	pyspark	idle	<a href="#">Link</a>	<a href="#">Link</a>	None	✓

SparkSession available as 'spark'.

```
In [2]: spark
```

```
<pyspark.sql.session.SparkSession object at 0xfffffa58c2730>
```

```
In [*]: # Load csv Dataset
# desde gdvie
#df=spark.read.csv('gdrive/MyDrive/st0263-252/bigdata/datasets/sample_data.csv',inferSchema=True,header=True)

# desde Local
#df=spark.read.csv('../datasets/sample_data.csv',inferSchema=True,header=True)
# desde S3
df=spark.read.csv('s3://emrsamuel/datasets/sample_data.csv',inferSchema=True,header=True)
```

```
In [*]: #columns of dataframe
df.columns
```

```
In [*]: #check number of columns
len(df.columns)
```

```
In [*]: #number of records in dataframe
df.count()
```

```
In [*]: #shape of dataset
print((df.count(),len(df.columns)))
```

```
In [*]: #printSchema
df.printSchema()
```

```
In [ ]: #fisrt few rows of dataframe
df.show(5)
```

```
In [ ]: #select only 2 columns
df.select('age','mobile').show(5)
```

jupyterhub Data\_processing\_using\_PySpark Last Checkpoint hace 3 minutos (autosaved)

Logout Control Panel

File Edit View Insert Cell Kernel Widgets Help Trusted PySpark O

```
df=spark.read.csv('s3://emrsamuel/datasets/sample_data.csv',inferSchema=True,header=True)
```

In [3]: #columns of dataframe  
df.columns  
['ratings', 'age', 'experience', 'family', 'mobile']

In [4]: #check number of columns  
len(df.columns)  
5

In [5]: #number of records in dataframe  
df.count()  
33

In [6]: #shape of dataset  
print((df.count(),len(df.columns)))  
(33, 5)

In [7]: #printSchema  
df.printSchema()

```
root
 |-- ratings: integer (nullable = true)
 |-- age: integer (nullable = true)
 |-- experience: double (nullable = true)
 |-- family: integer (nullable = true)
 |-- mobile: string (nullable = true)
```

In [8]: #first few rows of dataframe  
df.show(5)

	ratings	age	experience	family	mobile
0	3	32	9.0	3	Vivo
1	3	27	13.0	3	Apple
2	4	22	2.5	0	Samsung
3	4	37	16.5	4	Apple
4	5	27	9.0	1	MI

only showing top 5 rows

In [9]: #select only 2 columns

```
In [10]: #info about dataframe
df.describe().show()

+-----+-----+-----+-----+
|summary|      ratings|       age|   experience|     family|mobile|
+-----+-----+-----+-----+
| count|        33|        33|         33|        33|     33|
| mean| 3.57575757575757| 30.4848484848484| 10.3030303030303| 1.81818181818181|  NULL|
| stddev| 1.188806636071336|  6.18527087180309|  6.770731351213326| 1.8448330794164254|  NULL|
| min|        1|        22|         2.5|        0| Apple|
| max|        5|        42|        23.0|        5| Vivo|
+-----+-----+-----+-----+-----+-----+-----+
```

```
In [11]: from pyspark.sql.types import StringType,DoubleType,IntegerType
```

```
In [12]: #add column
df.withColumn("age_after_10_yrs", (df["age"]+10)).show(10,False)

+-----+-----+-----+-----+
|ratings|experience|family|mobile |age_after_10_yrs|
+-----+-----+-----+-----+
| 3 | 32 | 9.0 | 3 |Vivo | 42 |
| 3 | 27 |13.0 | 3 |Apple | 37 |
| 4 | 22 | 2.5 | 0 |Samsung| 32 |
| 4 | 37 |16.5 | 4 |Apple | 47 |
| 5 | 27 | 9.0 | 1 |MI | 37 |
| 4 | 27 | 9.0 | 0 |Oppo | 37 |
| 5 | 37 |23.0 | 5 |Vivo | 47 |
| 5 | 37 |23.0 | 5 |Samsung| 47 |
| 3 | 22 | 2.5 | 0 |Apple | 32 |
| 3 | 27 | 6.0 | 0 |MI | 37 |
+-----+-----+-----+-----+
only showing top 10 rows
```

```
In [13]: #add column
df.withColumn('age_double', df['age'].cast(DoubleType())).show(10,False)

+-----+-----+-----+-----+
|ratings|experience|family|mobile |age_double|
+-----+-----+-----+-----+
| 3 | 32 | 9.0 | 3 |Vivo | 32.0 |
| 3 | 27 |13.0 | 3 |Apple | 27.0 |
| 4 | 22 | 2.5 | 0 |Samsung| 22.0 |
| 4 | 37 |16.5 | 4 |Apple | 37.0 |
| 5 | 27 | 9.0 | 1 |MI | 27.0 |
| 4 | 27 | 9.0 | 0 |Oppo | 27.0 |
| 5 | 37 |23.0 | 5 |Vivo | 37.0 |
+-----+-----+-----+-----+
```

jupyterhub Data\_processing\_using\_PySpark Last Checkpoint: hace 3 minutos (autosaved) Logout Control Panel

File Edit View Insert Cell Kernel Widgets Help Trusted PySpark

In [23]: # Value counts  
df.groupBy('mobile').mean().show(5,False)

mobile	avg(ratings)	avg(age)	avg(experience)	avg(family)
Samsung	4.1666666666666667	28.666666666666668	8.666666666666666	1.8333333333333333
MI	3.5	30.125	10.1875	1.375
Oppo	2.857142857142857	28.4285714285714271	10.357142857142858	1.4285714285714286
Apple	3.4285714285714284	38.571428571428573	11.0	2.7142857142857144
Vivo	4.2	36.0	11.4	1.8

In [24]: df.groupBy('mobile').sum().show(5,False)

mobile	sum(ratings)	sum(age)	sum(experience)	sum(family)
Samsung	25	172	52.0	11
MI	28	241	81.5	11
Oppo	20	199	72.5	10
Apple	24	214	77.0	19
Vivo	21	180	57.0	9

In [25]: # Value counts  
df.groupBy('mobile').max().show(5,False)

mobile	max(ratings)	max(age)	max(experience)	max(family)
Samsung	5	37	23.0	5
MI	5	42	23.0	5
Oppo	4	42	23.0	2
Apple	4	37	16.5	5
Vivo	5	37	23.0	5

In [26]: # Value counts  
df.groupBy('mobile').min().show(5,False)

mobile	min(ratings)	min(age)	min(experience)	min(family)
Samsung	2	22	2.5	0

jupyterhub Data\_processing\_using\_PySpark Last Checkpoint: hace 3 minutos (autosaved)

Logout Control Panel

File Edit View Insert Cell Kernel Widgets Help Trusted PySpark

```
In [38]: #duplicate values
df.count()

33

In [39]: #drop duplicate values
df=df.dropDuplicates()

In [40]: #validate new count
df.count()

26

In [41]: #drop column of dataframe
df_new=df.drop('mobile')

In [42]: df_new.show(10)

+---+---+---+
|ratings|age|experience|family|
+---+---+---+
| 4| 22| 2.5| 0|
| 4| 22| 6.0| 1|
| 3| 27| 6.0| 0|
| 2| 32| 16.5| 2|
| 4| 27| 9.0| 0|
| 3| 37| 16.5| 5|
| 4| 27| 6.0| 1|
| 4| 37| 9.0| 2|
| 3| 22| 2.5| 0|
| 3| 32| 9.0| 3|
+---+---+---+
only showing top 10 rows

In [ ]: # saving file to csv

In [ ]: #current working directory
!pwd

In [49]: #target directory
#pathcsv_out='./out/df_csv'
#pathcsv_out='gdrive/MyDrive/st0263-252/out/df_csv'
# hacia S3
write_uri = 's3://emrsamuel/df_csv'
```

jupyterhub Data\_processing\_using\_PySpark Last Checkpoint hace 4 minutos (autosaved)

Logout Control Panel

File Edit View Insert Cell Kernel Widgets Help Trusted PySpark

```
In [42]: df_new.show(10)
+-----+-----+-----+
|ratings|age|experience|family|
+-----+-----+-----+
| 4| 22| 2.5| 0|
| 4| 22| 6.0| 1|
| 3| 27| 6.0| 0|
| 2| 32| 16.5| 2|
| 4| 27| 9.0| 0|
| 3| 37| 16.5| 5|
| 4| 27| 6.0| 1|
| 4| 37| 9.0| 2|
| 3| 22| 2.5| 0|
| 3| 32| 9.0| 3|
+-----+-----+
only showing top 10 rows
```

```
In [ ]: # saving file to csv
```

```
In [ ]: #current working directory
!pwd
```

```
In [49]: #target directory
#pathcsv_out='./out/df_csv'
#pathcsv_out='gdrive/MyDrive/st0263-252/out/df_csv'
# hacia S3
write_uri = 's3://emrsamuel/df_csv'
```

```
In [51]: #save the dataframe as single csv
df.coalesce(1).write.format("csv").option("header","true").save(write_uri)
```

```
In [ ]: # parquet
```

```
In [52]: #target location
#pathparquet_out='./out/df_parquet'
#pathparquet_out='gdrive/MyDrive/st0263-252/out/df_parquet'

# hacia S3
write_uri='s3://emrsamuel/df_parquet'
```

```
In [54]: #save the data into parquet format
df.write.mode("overwrite").format("parquet").save(write_uri)
```

emrsamuel > df\_csv/

**df\_csv/**

Objetos Propiedades

**Objetos (2)**

Los objetos son las entidades fundamentales que se almacenan en Amazon S3. Puede utilizar el [índice de Amazon S3](#) para obtener una lista de todos los objetos de su bucket. Para que otras personas obtengan acceso que concederles permisos de forma explícita. [Más información](#)

Copiar URI de S3 Copiar URL Descargar Abrir i Eliminar Acciones Crear carpeta

Nombre	Tipo	Última modificación	Tamaño	Clase de almacenamiento
_SUCCESS	-	17 Nov 2025 12:06:34 AM -05	0 B	Estándar
part-00000-94c7f10f-b96b-4e8b-97a6-055176b5a4f6-c000.csv	csv	17 Nov 2025 12:06:34 AM -05	474.0 B	Estándar

> emrsamuel > df\_parquet/

df\_parquet/

**Objetos** Propiedades

**Objetos (2)**

Los objetos son las entidades fundamentales que se almacenan en Amazon S3. Puede utilizar el [inventario de Amazon S3](#) para obtener una lista de todos los objetos de su bucket. Para que otras personas lo que concederles permisos de forma explícita. [Más información](#)

Buscar objetos por prefijo:

Nombre	Tipo	Última modificación	Tamaño	Clase de
_SUCCESS	-	17 Nov 2025 12:07:21 AM -05	0 B	Estándar
part-00000-8420689f-6bd4-4bc8-adcb-ac20dd182c98-c000.snappy.parquet	parquet	17 Nov 2025 12:07:20 AM -05	1.7 KB	Estándar

jupyterhub spark\_colab\_ejercicios Last Checkpoint: hace 16 minutos (unsaved changes)

Logout Control Panel

File Edit View Insert Cell Kernel Widgets Help Trusted PySpark O

Ejemplo 1: WordCount con RDD

```
In [1]: from pyspark import SparkContext
sc = SparkContext.getOrCreate()

text = sc.textFile('s3://emrsamuel/datasets/gutenberg/gutenberg-txt-es.zip-url.txt')
# Simular archivo de texto
# text = sc.parallelize(["Hola Spark Hola Big Data", "Spark es rápido y poderoso"])
counts = text.flatMap(lambda x: x.split(" ")) \
    .map(lambda x: (x, 1)) \
    .reduceByKey(lambda a, b: a + b)
counts.collect()
```

Starting Spark application

ID	YARN Application ID	Kind	State	Spark UI	Driver log	User	Current session?
7	application_1763344109923_0010	pyspark	idle	<a href="#">Link</a>	<a href="#">Link</a>	None	✓

SparkSession available as 'spark'.

```
[('https://drive.google.com/open?id=1jz-p_5gP7TVxhJMrx8eO1J2ZqsSupR_w', 1), ('', 1)]
```

Ejemplo 2: Análisis con DataFrame API

```
In [2]: from pyspark.sql import SparkSession
spark = SparkSession.builder.getOrCreate()

# Simular DataFrame de ventas
data = [("martillo", 12000), ("taladro", 45000), ("martillo", 15000)]
columns = ["producto", "valor"]
df = spark.createDataFrame(data, columns)
df.groupBy("producto").sum("valor").show()
```

producto	sum(valor)
martillo	27000
taladro	45000

No es seguro [https://ec2-3-236-117-255.compute-1.amazonaws.com:9443/user/jovyan/notebooks/spark\\_colab\\_ejercicios.ipynb#](https://ec2-3-236-117-255.compute-1.amazonaws.com:9443/user/jovyan/notebooks/spark_colab_ejercicios.ipynb#)

jupyterhub spark\_colab\_ejercicios Last Checkpoint: hace 16 minutos (unsaved changes)

File Edit View Insert Cell Kernel Widgets Help Logout Control Panel Trusted PySpark

TASK Application ID KILLED State SPARK UI Driver Log User Current Session:

7 application\_1763344109923\_0010 pyspark idle Link Link None ✓

SparkSession available as 'spark'.

[('https://drive.google.com/open?id=1jz-p\_5gP7TVxhJMrx8eO1J2ZqsSupR\_w', 1), ('', 1)]

### Ejemplo 2: Análisis con DataFrame API

```
In [2]: from pyspark.sql import SparkSession
spark = SparkSession.builder.getOrCreate()

# Simular DataFrame de ventas
data = [("martillo", 12000), ("taladro", 45000), ("martillo", 15000)]
columns = ["producto", "valor"]
df = spark.createDataFrame(data, columns)
df.groupBy("producto").sum("valor").show()
```

producto	sum(valor)
martillo	27000
taladro	45000

### Ejemplo 3: Clasificación con MLlib

```
In [4]: from pyspark.ml.feature import VectorAssembler
from pyspark.ml.classification import LogisticRegression

df = spark.read.csv("s3://emrsamuel/datasets/clientes.csv", header=True, inferSchema=True)

assembler = VectorAssembler(inputCols=["edad", "ingresos"], outputCol="features")
data = assembler.transform(df).select("features", df["comprador"].alias("label"))
train, test = data.randomSplit([0.8, 0.2], seed=42)
lr = LogisticRegression()
model = lr.fit(train)
model.transform(test).select("features", "label", "prediction").show()
```

features	label	prediction
[34.0,4500.0]	1	0.0

jupyterhub spark\_colab\_ejercicios Last Checkpoint: hace 31 minutos (autosaved)

File Edit View Insert Cell Kernel Widgets Help

Logout Control Panel

Not Trusted PySpark

In [1]:

```
from pyspark import SparkContext
sc = SparkContext.getOrCreate()

text = sc.textFile('s3://emrsamuel/datasets/gutenberg/gutenberg-txt-es.zip-url.txt')
# Simular archivo de texto
# text = sc.parallelize(["Hola Spark Hola Big Data", "Spark es rápido y poderoso"])
counts = text.flatMap(lambda x: x.split(" ")) \
    .map(lambda x: (x, 1)) \
    .reduceByKey(lambda a, b: a + b)
counts.collect()

Starting Spark application

ID          YARN Application ID   Kind  State  Spark UI  Driver log  User  Current session?
10  application_1763344109923_0013  pyspark  idle           None      ✓

SparkSession available as 'spark'.
[('https://drive.google.com/open?id=1jz-p_5gP7TVxhJMrx8e01J2ZqsSupR_w', 1), ('', 1)]
```

**Ejemplo 2: Análisis con DataFrame API**

In [2]:

```
from pyspark.sql import SparkSession
spark = SparkSession.builder.getOrCreate()

# Simular DataFrame de ventas
data = [("martillo", 12000), ("taladro", 45000), ("martillo", 15000)]
columns = ["producto", "valor"]
df = spark.createDataFrame(data, columns)
df.groupBy("producto").sum("valor").show()
```

producto	sum(valor)
martillo	27000
taladro	45000

**Ejemplo 3: Clasificación con MLlib**

In [3]:

```
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.classification import LogisticRegression

df = spark.read.csv("s3://emrsamuel/datasets/clientes.csv", header=True, inferSchema=True)
```

No es seguro [https://ec2-3-236-117-255.compute-1.amazonaws.com:9443/user/jovyan/notebooks/spark\\_colab\\_ejercicios.ipynb#](https://ec2-3-236-117-255.compute-1.amazonaws.com:9443/user/jovyan/notebooks/spark_colab_ejercicios.ipynb#)

jupyterhub spark\_colab\_ejercicios Last Checkpoint: hace 31 minutos (autosaved)

File Edit View Insert Cell Kernel Widgets Help

Logout Control Panel Not Trusted PySpark O

```
10 application_1763344109923_0013 pyspark idle None ✓
SparkSession available as 'spark'.
[('https://drive.google.com/open?id=1jz-p_5gP7TVxhJMrx8e01J2ZqsSupR_w', 1), ('', 1)]
```

### Ejemplo 2: Análisis con DataFrame API

```
In [2]: from pyspark.sql import SparkSession
spark = SparkSession.builder.getOrCreate()

# Simular DataFrame de ventas
data = [("martillo", 12000), ("taladro", 45000), ("martillo", 15000)]
columns = ["producto", "valor"]
df = spark.createDataFrame(data, columns)
df.groupBy("producto").sum("valor").show()
```

```
+-----+-----+
|producto|sum(valor)|
+-----+-----+
|martillo|    27000|
| taladro|    45000|
+-----+-----+
```

### Ejemplo 3: Clasificación con MLlib

```
In [3]: from pyspark.ml.feature import VectorAssembler
from pyspark.ml.classification import LogisticRegression

df = spark.read.csv("s3://emrsamuel/datasets/clientes.csv", header=True, inferSchema=True)

assembler = VectorAssembler(inputCols=["edad", "ingresos"], outputCol="features")
data = assembler.transform(df).select("features", df["comprador"].alias("label"))
train, test = data.randomSplit([0.8, 0.2], seed=42)
lr = LogisticRegression()
model = lr.fit(train)
model.transform(test).select("features", "label", "prediction").show()
```

```
+-----+-----+-----+
|     features|label|prediction|
+-----+-----+-----+
|[34.0,4500.0]|   1|      0.0|
+-----+-----+-----+
```

resultados/

Objetos Propiedades

Objetos (8)

Los objetos son las entidades fundamentales que se almacenan en Amazon S3. Puede utilizar el inventario de Amazon S3 para obtener una lista de todos los objetos de su bucket. Para que otras personas obtengan acceso a sus objetos, tendrá que concederles permisos de forma explícita. [Más información](#)

Copiar URI de S3 Copiar URL Descargar Abrir Eliminar Acciones Crear carpeta Cargar

Nombre	Tipo	Última modificación	Tamaño	Clase de almacenamiento
<a href="#">ciudades.csv/</a>	Carpetas	-	-	-
<a href="#">ciudades.parquet/</a>	Carpetas	-	-	-
<a href="#">departamentos.csv/</a>	Carpetas	-	-	-
<a href="#">departamentos.parquet/</a>	Carpetas	-	-	-
<a href="#">edades.csv/</a>	Carpetas	-	-	-
<a href="#">edades.parquet/</a>	Carpetas	-	-	-
<a href="#">fechas.csv/</a>	Carpetas	-	-	-
<a href="#">fechas.parquet/</a>	Carpetas	-	-	-

## Ejercicio adicional propuesto en el Readme

jupyterhub Untitled4 Last Checkpoint: hace 9 minutos (autosaved)

File Edit View Insert Cell Kernel Help

Logout Control Panel Trusted PySpark

only showing top 5 rows

In [13]: df.printSchema()

```
root
|-- fecha_reporte_web: string (nullable = true)
|-- ID de caso: integer (nullable = true)
|-- Fecha de notificación: string (nullable = true)
|-- Código DIVIPOLA departamento: integer (nullable = true)
|-- Nombre departamento: string (nullable = true)
|-- Código DIVIPOLA municipio: integer (nullable = true)
|-- Nombre municipio: string (nullable = true)
|-- Edad: integer (nullable = true)
|-- Unidad de medida de edad: integer (nullable = true)
|-- Sexo: string (nullable = true)
|-- Tipo de contagio: string (nullable = true)
|-- Ubicación del caso: string (nullable = true)
|-- Estado: string (nullable = true)
|-- Código ISO del país: integer (nullable = true)
|-- Nombre del país: string (nullable = true)
|-- Recuperado: string (nullable = true)
|-- Fecha de inicio de síntomas: string (nullable = true)
|-- Fecha de muerte: string (nullable = true)
|-- Fecha de diagnóstico: string (nullable = true)
|-- Fecha de recuperación: string (nullable = true)
|-- Tipo de recuperación: string (nullable = true)
|-- Pertenencia étnica: integer (nullable = true)
|-- Nombre del grupo étnico: string (nullable = true)
```

In [14]: df = df.toDF(\*[c.strip().lower().replace(" ", "\_").replace("ó","o").replace("á","a").replace("í","i").replace("é","e").replace("í","i").replace("ñ","n") for c in df.columns]).show(5)

```
+-----+-----+-----+-----+-----+
|fecha_reporte_web|id_de_caso|fecha_de_notificación|código_divipola_departamento|nombre_departamento|
|nombre_municipio|edad|unidad_de_medida_de_edad|sexo|tipo_de_contagio|ubicación_del_caso|estado|código_iso_del_país|nombre_del_p
aís|recuperado|fecha_de_inicio_de_síntomas|fecha_de_muerte|fecha_de_diagnóstico|fecha_de_recuperación|tipo_de_recuperación|pert
enencia étnica|nombre_del_grupo étnico|
+-----+-----+-----+-----+-----+
```

jupyterhub Untitled4 Last Checkpoint: hace 9 minutos (autosaved)

File Edit View Insert Cell Kernel Widgets Help Trusted PySpark

```
df = df.withColumn("edad_grupo", when((col("edad") < 18), "menor_18").when((col("edad") >= 18) & (col("edad") <= 40), "18-40").when((col("edad") > 40) & (col("edad") <= 60), "41-60").otherwise("mayor_60"))  
df.select("edad", "edad_grupo").show(10)  
  
+---+-----+  
|edad|edad_grupo|  
+---+-----+  
| 19| 18-40|  
| 34| 18-40|  
| 50| 41-60|  
| 55| 41-60|  
| 25| 18-40|  
| 27| 18-40|  
| 85| mayor_60|  
| 22| 18-40|  
| 28| 18-40|  
| 36| 18-40|  
+---+-----+  
only showing top 10 rows  
  
In [16]: df_activos = df.filter(col("estado") == "Leve") # o el valor que indique activo  
df_activos.show(5)  
  
+-----+-----+-----+-----+-----+  
+-----+-----+-----+-----+-----+  
+-----+-----+-----+-----+-----+  
|fecha_reporte_web|id_de_caso|fecha_de_notificación|código_divipola_departamento|nombre_departamento|código_divipola_municipio|  
|nombre_municipio|edad|unidad_de_medida_de_edad|sexo|tipo_de_contagio|ubicación_del_caso|estado|código_iso_del_país|nombre_del_p  
|aís|recuperado|fecha_de_inicio_de_síntomas|fecha_de_muerte|fecha_de_diagnóstico|fecha_de_recuperación|tipo_de_recuperación|pert  
enencia étnica|nombre_del_grupo étnico|edad_grupo|  
+-----+-----+-----+-----+-----+  
+-----+-----+-----+-----+-----+  
+-----+-----+-----+-----+-----+  
+-----+-----+-----+-----+-----+  
| 6/3/2020 0:00:00| 1| 2/3/2020 0:00:00| 11| BOGOTÁ| 11001|  
|BOGOTÁ| 19| F| Importado| Casa| Leve| 380| ITALIA| Recupe  
|rado| 27/2/2020 0:00:00| NULL| 6/3/2020 0:00:00| 13/3/2020 0:00:00| PCR|  
| 6| NULL| 18-40|  
| 9/3/2020 0:00:00| 2| 6/3/2020 0:00:00| 76| VALLE| 76111|
```

File Edit View Insert Cell Kernel Widgets Help

Trusted PySpark

```
In [17]: df_departamento = df.groupBy("nombre_departamento")\
    .count()\
    .orderBy("count", ascending=False)
df_departamento.show(10)
```

nombre_departamento	count
BOGOTA	401
VALLE	148
ANTIOQUIA	106
CUNDINAMARCA	49
CARTAGENA	39
RISARALDA	34
BARRANQUILLA	31
HUILA	30
QUINDIO	23
NORTE SANTANDER	19

only showing top 10 rows

```
In [18]: df_ciudad = df.groupBy("nombre_municipio")\
    .count()\
    .orderBy("count", ascending=False)
df_ciudad.show(10)
```

nombre_municipio	count
BOGOTA	401
CALI	101
MEDELLIN	63
CARTAGENA	39
BARRANQUILLA	31
NEIVA	27
PEREIRA	25
PALMIRA	22
VALLEDUPAR	16
CUCUTA	15

only showing top 10 rows

```
In [19]: df_fecha = df.groupBy("fecha_reporte_web")\
    .count()\
    .orderBy("count", ascending=False)
df_fecha.show(10)
```

jupyterhub Untitled4 Last Checkpoint: hace 9 minutos (autosaved)

Logout Control Panel

File Edit View Insert Cell Kernel Widgets Help Trusted PySpark

In [20]:

```
df_edad = df.groupBy("edad_grupo")\
    .count()\
    .orderBy("edad_grupo")
df_edad.show()
```

edad_grupo	count
18-40	476
41-60	332
mayor_60	159
menor_18	33

In [21]:

```
df.createOrReplaceTempView("covid")
```

In [22]:

```
spark.sql("""
SELECT nombre_departamento, COUNT(*) AS casos
FROM covid
GROUP BY nombre_departamento
ORDER BY casos DESC
LIMIT 10
""").show()
```

nombre_departamento	casos
BOGOTA	401
VALLE	148
ANTIOQUIA	106
CUNDINAMARCA	49
CARTAGENA	39
RISARALDA	34
BARRANQUILLA	31
HUILA	30
QUINDIO	23
NORTE SANTANDER	19

In [23]:

```
spark.sql("""
SELECT nombre_municipio, COUNT(*) AS casos
FROM covid
GROUP BY nombre_municipio
```

```

In [24]: spark.sql("""
    SELECT fecha_reporte_web, COUNT(*) AS casos
    FROM covid
    GROUP BY fecha_reporte_web
    ORDER BY casos DESC
    LIMIT 10
""").show()

+-----+-----+
|fecha_reporte_web|casos|
+-----+-----+
|31/3/2020 0:00:00| 107|
|24/3/2020 0:00:00| 105|
|30/3/2020 0:00:00| 96|
|1/4/2020 0:00:00| 95|
|29/3/2020 0:00:00| 94|
|23/3/2020 0:00:00| 73|
|28/3/2020 0:00:00| 69|
|25/3/2020 0:00:00| 61|
|27/3/2020 0:00:00| 49|
|20/3/2020 0:00:00| 48|
+-----+-----+


In [25]: spark.sql("""
    SELECT
        CASE
            WHEN edad < 18 THEN 'menor_18'
            WHEN edad BETWEEN 18 AND 40 THEN '18-40'
            WHEN edad BETWEEN 41 AND 60 THEN '41-60'
            ELSE 'mayor_60'
        END AS edad_grupo,
        COUNT(*) AS casos
    FROM covid
    GROUP BY
        CASE
            WHEN edad < 18 THEN 'menor_18'
            WHEN edad BETWEEN 18 AND 40 THEN '18-40'
            WHEN edad BETWEEN 41 AND 60 THEN '41-60'
            ELSE 'mayor_60'
        END
    ORDER BY edad_grupo
""").show()

+-----+-----+
|edad_grupo|casos|
+-----+-----+

```

## Ejecución de .py en main node

```

25/11/17 06:24:13 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-Java classes where applicable
25/11/17 06:24:14 INFO DefaultNoHDFSFailoverProxyProvider: Connecting to ResourceManager at ip-172-31-71-159.ec2.internal/172.31.71.159:8082
25/11/17 06:24:14 INFO Configuration: resource-types.xml not found
25/11/17 06:24:14 INFO ResourceUtils: Unable to find 'resource-types.xml'.
25/11/17 06:24:14 INFO Client: Verifying our application has requested no more than the maximum memory capability of the cluster (23424 MB per container)
25/11/17 06:24:14 INFO Client: Submitting application with 2423 MB memory including 384 MB overhead
25/11/17 06:24:14 INFO Client: Setting up container launch context for our AM
25/11/17 06:24:14 INFO Client: Setting up the launch environment for our AM container
25/11/17 06:24:14 WARN Client: Neither spark.yarn.jars nor spark.yarn.archive is set, falling back to uploading libraries under SPARK_HOME.
25/11/17 06:24:15 INFO Client: Uploading resource file:/mnt/tmp/spark-aef0f01a1-lala-4811-997f-a06ea3808714/_spark_libs_1722523502337600274.zip -> hdfs://ip-172-31-71-159.ec2.internal:8020/user/hadoop/.sparkStaging/application_1763344109923_0020/pyspark.zip
25/11/17 06:24:15 INFO Client: Uploading resource file:/etc/spark/conf.dist/hive-site.xml -> hdfs://ip-172-31-71-159.ec2.internal:8020/user/hadoop/.sparkStaging/application_1763344109923_0020/hive-site.xml
25/11/17 06:24:15 INFO Client: Uploading resource file:/etc/spark/conf.dist/hudi-defaults.conf -> hdfs://ip-172-31-71-159.ec2.internal:8020/user/hadoop/.sparkStaging/application_1763344109923_0020/hudi-defaults.conf
25/11/17 06:24:16 INFO Client: Uploading resource file:/home/hadoop/stb263-252/bigdata/03-spark/wc-pyspark.py -> hdfs://ip-172-31-71-159.ec2.internal:8020/user/hadoop/.sparkStaging/application_1763344109923_0020/wc-pyspark.py
25/11/17 06:24:16 INFO Client: Uploading resource file:/usr/lib/spark/python/lib/pySpark.zip -> hdfs://ip-172-31-71-159.ec2.internal:8020/user/hadoop/.sparkStaging/application_1763344109923_0020/pySpark.zip
25/11/17 06:24:16 INFO Client: Uploading resource file:/usr/lib/spark/python/lib/py4j-0.10.9.7-src.zip -> hdfs://ip-172-31-71-159.ec2.internal:8020/user/hadoop/.sparkStaging/application_1763344109923_0020/py4j-0.10.9.7-src.zip
25/11/17 06:24:16 INFO SecurityManager: Changing view acls to: hadoop
25/11/17 06:24:16 INFO SecurityManager: Changing view acls groups to:
25/11/17 06:24:16 INFO SecurityManager: Changing modify acls to: hadoop
25/11/17 06:24:16 INFO SecurityManager: Changing modify acls groups to:
25/11/17 06:24:16 INFO SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users with view permissions: hadoop; groups with view permissions: EMPTY; users with modify permissions: hadoop; groups with modify permissions: hadoop
25/11/17 06:24:16 INFO Client: Submitted application application_1763344109923_0020
25/11/17 06:24:17 INFO Client: Application report for application_1763344109923_0020 (state: ACCEPTED)
25/11/17 06:24:17 INFO Client:
    client token: N/A
    diagnostics: AM container is launched, waiting for AM container to Register with RM
    ApplicationMaster host: N/A
    ApplicationMaster RPC port: -1
    queue: root.default
    start time: 1763360656190
    final status: UNDEFINED
    tracking URL: http://ip-172-31-71-159.ec2.internal:20888/proxy/application_1763344109923_0020/
user: hadoop
25/11/17 06:24:20 INFO Client: Application report for application_1763344109923_0020 (state: RUNNING)
25/11/17 06:24:20 INFO Client:
    client token: N/A
    diagnostics: N/A
    ApplicationMaster host: ip-172-31-70-28.ec2.internal
    ApplicationMaster RPC port: 37287
    queue: root.default
    start time: 1763360656190
    final status: UNDEFINED
    tracking URL: http://ip-172-31-71-159.ec2.internal:20888/proxy/application_1763344109923_0020/
user: hadoop

```

aws | [Alt+S]

Amazon S3 > Buckets > emrsamuel > datasets/ > gutenberg/

## CloudShell

us-east-1 +

```
25/11/17 06:24:14 INFO DefaultNoHARFailoverProxyProvider: Connecting to ResourceManager at ip-172-31-71-159.ec2.internal/172.31.71.159:8032
25/11/17 06:24:14 INFO Configuration: resource-types.xml not found
25/11/17 06:24:14 INFO ResourceUtils: Unable to find 'resource-types.xml'.
25/11/17 06:24:14 INFO Client: Verifying our application has not requested more than the maximum memory capability of the cluster (23424 MB per container).
25/11/17 06:24:14 INFO Client: Will allocate AM container, with 2432 MB memory including 384 MB overhead
25/11/17 06:24:14 INFO Client: Setting up container launch context for our AM
25/11/17 06:24:14 INFO Client: Setting up the launch environment for our AM container
25/11/17 06:24:14 INFO Client: Preparing resources for our AM container
25/11/17 06:24:14 WARN Client: Neither spark.yarn.jars nor spark.yarn.archive is set, falling back to uploading libraries under SPARK_HOME.
25/11/17 06:24:15 INFO Client: Uploading resource file:/mnt/tmp/spark-ae0f01a1-1ala-4811-997f-a06ea3808714/_spark_libs_13722523502337600274.zip -00274.zip
25/11/17 06:24:15 INFO Client: Uploading resource file:/etc/spark/conf.dist/hive-site.xml -> hdfs://ip-172-31-71-159.ec2.internal:8020/user/hadoop/.hive-site.xml
25/11/17 06:24:16 INFO Client: Uploading resource file:/etc/hudi/conf.dist/hudi-defaults.conf -> hdfs://ip-172-31-71-159.ec2.internal:8020/user/hadoop/.hudi-defaults.conf
25/11/17 06:24:16 INFO Client: Uploading resource file:/home/hadoop/st0263-252/bigdata/03-spark/wc-pyspark.py -> hdfs://ip-172-31-71-159.ec2.internal:8020/user/hadoop/wc-pyspark.py
25/11/17 06:24:16 INFO Client: Uploading resource file:/usr/lib/spark/python/lib/pyspark.zip -> hdfs://ip-172-31-71-159.ec2.internal:8020/user/hadoop/pyspark.zip
25/11/17 06:24:16 INFO Client: Uploading resource file:/usr/lib/spark/python/lib/py4j-0.10.9.7-src.zip -> hdfs://ip-172-31-71-159.ec2.internal:8020/user/hadoop/py4j-0.10.9.7-src.zip
25/11/17 06:24:16 INFO Client: Uploading resource file:/mnt/tmp/spark-ae0f01a1-1ala-4811-997f-a06ea3808714/_spark_conf_8726557474209513041.zip -0020.zip
25/11/17 06:24:16 INFO SecurityManager: Changing view acls to: hadoop
25/11/17 06:24:16 INFO SecurityManager: Changing modify acls to: hadoop
25/11/17 06:24:16 INFO SecurityManager: Changing view acls groups to:
25/11/17 06:24:16 INFO SecurityManager: Changing modify acls groups to:
25/11/17 06:24:16 INFO SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users with view permissions: hadoop; groups with view permissions: 
25/11/17 06:24:16 INFO Client: Submitting application application_1763344109923_0020 to ResourceManager
25/11/17 06:24:16 INFO YarnClientImpl: Submitted application application_1763344109923_0020
25/11/17 06:24:17 INFO Client: Application report for application_1763344109923_0020 (state: ACCEPTED)
25/11/17 06:24:17 INFO Client:
  client token: N/A
  diagnostics: AM container is launched, waiting for AM container to Register with RM
  ApplicationMaster host: N/A
  ApplicationMaster RPC port: -1
  queue: root.default
  start time: 1763344109923
  final status: UNDEFINED
  tracking URL: http://ip-172-31-71-159.ec2.internal:20888/proxy/application_1763344109923_0020/
  user: hadoop
25/11/17 06:24:20 INFO Client: Application report for application_1763344109923_0020 (state: RUNNING)
25/11/17 06:24:20 INFO Client:
  client token: N/A
  diagnostics: N/A
  ApplicationMaster host: ip-172-31-70-28.ec2.internal
  ApplicationMaster RPC port: 37287
  queue: root.default
  start time: 1763344109923
  final status: UNDEFINED
  tracking URL: http://ip-172-31-71-159.ec2.internal:20888/proxy/application_1763344109923_0020/
  user: hadoop
25/11/17 06:24:50 INFO Client: Application report for application_1763344109923_0020 (state: RUNNING)
```

```
#configuración de Spark en Google Colab
import os
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-17-openjdk-amd64"
os.environ["SPARK_HOME"] = "/content/spark-4.0.1-bin-hadoop3"

#configuración de Spark en Google Colab
import findspark
findspark.init()

#configuración de Spark en Google Colab
from pyspark.sql import SparkSession
spark = SparkSession.builder.master("local[*]").getOrCreate()
sc = spark.sparkContext

# WORDCOUNT COMPACTO
# en AWS S3
#files_rdd = sc.textFile("s3a://username_datalake/datasets/gutenberg-small/*.txt")

# en gdrive:
files_rdd = sc.textFile("gdrive/MyDrive/st0263-252/bigdata/datasets/gutenberg-small/*.txt")

# local:
#files_rdd = sc.textFile("../datasets/gutenberg-small/*.txt")
wc_unsort = files_rdd.flatMap(lambda line: line.split()).map(lambda word: (word, 1)).reduceByKey(lambda a, b: a + b)
wc = wc_unsort.sortBy(lambda a: -a[1])
for tupla in wc.take(10):
    print(tupla)

('the', 44647)
('of', 28020)
('to', 23208)
('and', 20444)
('in', 13174)
('that', 12265)
('I', 10880)
('a', 10431)
('is', 7776)
('be', 7148)

# WORDCOUNT PASO A PASO
```

The screenshot shows a Jupyter Notebook interface with the following details:

- Title Bar:** wordcount-spark-colab.ipynb
- Menu Bar:** Archivo, Editar, Ver, Insertar, Entorno de ejecución, Herramientas, Ayuda
- Toolbar:** Comandos, Código, Texto, Ejecutar todas
- Code Cells:**
  - [9] 0 s: The Biblophile Society
  - [10] 5 s: A code cell showing the creation of a word count RDD `wc1` by mapping word counts from tokens, followed by printing the top 10 results. The output includes words like 'LINCOLN', 'LETTERS', 'By', 'Abraham', 'Lincoln', 'Published', 'by', 'The', 'Biblophile', and 'Society'.
  - [10] 5 s: Reduces the word counts and prints the top 10 results again. The output includes 'Published', 'themselves', 'were', 'sheet', 'despatched', 'most', 'turbulent', 'A.', 'ORIGINALS', and 'IN'.
  - [11] 12 s: Sorts the word counts in descending order and prints the top 10 results. The output includes 'the', 'of', 'to', 'and', 'in', 'that', 'I', 'a', 'is', and 'be'.
- Variables and Terminal:** Buttons for Variables and Terminal are visible at the bottom left.
- Decorations:** A blue circular icon with a white arrow is located in the bottom right corner.

```
> ⌂ colab.research.google.com/drive/1btERLGa9OjH7B62J8y8bRwGa9Ryd5P_m#scrollTo=66f6a00c ⌂ 🔍 ⌂ 🎨

Comandos + Código + Texto | Ejecutar todas ▾

[3] ✓ 9s
from pyspark import SparkContext
sc = SparkContext.getOrCreate()

text = sc.textFile("gdrive/MyDrive/st0263-252/bigdata/datasets/gutenberg-small/*.txt")
# Simular archivo de texto
# text = sc.parallelize(["Hola Spark Hola Big Data", "Spark es rápido y poderoso"])
counts = text.flatMap(lambda x: x.split(" ")) \
    .map(lambda x: (x, 1)) \
    .reduceByKey(lambda a, b: a + b)
counts.collect()

('statutes', 5),
('Lord', 6),
('teeth', 15),
('displeased', 4),
('conditionally', 6),
('iti', 1),
('I', 14),
('leg', 2),
('frankly', 4),
('complimented', 6),
('hin'), 3),
('entertain', 2),
('prohibits', 4),
('Because', 13),
('citizenship...', 1),
('affirmation', 2),
('matter...', 1),
('Harris', 4),
('narrowed', 5),
('exonerate', 3),
('show', 7),
('Yates', 4),
("Harris's", 3),
('charity;', 1),
('enlarging', 15),
('resisting', 16),
('misplaced', 3),
('man,-I', 1),
('Democrat', 11),
('truth?', 5),
('Tract', 4),
('irresistible', 5),
('about,-a', 1),
('please...', 1),
('dressing', 1),
('positively', 13),
('sapper', 2),
('misunderstood', 9),
('incipient', 5),
('antagonism', 4),
('ignores', 3),
('Douglas;', 6),
('accordin', 1)
```

Comandos + Código + Texto Ejecutar todas

```
[4] ✓ 6s
dt = spark.createDataFrame(data, columns)
df.groupBy("producto").sum("valor").show()

+-----+-----+
|producto|sum(valor)|
+-----+-----+
|martillo|     27800|
|taladro|    45000|
+-----+-----+
```

Ejemplo 3: Clasificación con MLlib

```
[5] ✓ 7s
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.classification import LogisticRegression

df = spark.read.csv("gdrive/MyDrive/st0263-252/bigdata/datasets/clientes.csv", header=True, inferSchema=True)

assembler = VectorAssembler(inputCols=["edad", "ingresos"], outputCol="features")
data = assembler.transform(df).select("features", df["comprador"].alias("label"))
train, test = data.randomSplit([0.8, 0.2], seed=42)
lr = LogisticRegression()
model = lr.fit(train)
model.transform(test).select("features", "label", "prediction").show()

+-----+-----+-----+
| features|label|prediction|
+-----+-----+-----+
|[34.0,4500.0]|   1|       0.0|
+-----+-----+-----+
```

Ejemplo 4: Spark GraphX

```
[6] ✓ min
!pip install -q pyspark
!pyspark --packages graphframes:graphframes:0.8.3-spark3.5-s_2.12

Python 3.12.12 (main, Oct 10 2025, 08:52:57) [GCC 11.4.0] on linux
Type "help", "copyright", "credits" or "license" for more information.
WARNING: Using incubator modules: jdk.incubator.vector
:: loading settings :: url = jar:file:/content/spark-4.0.1-bin-hadoop3/jars/ivy-2.5.3.jar!/org/apache/ivy/core/settings/ivysettings.xml
Ivy Default Cache set to: /root/.ivy2.5.2/cache
The jars for the packages stored in: /root/.ivy2.5.2/jars
graphframes#graphframes added as a dependency
:: resolving dependencies :: org.apache.spark#spark-submit-parent-1704c169-6488-4bbe-bd1d-54a989d0920f;1.0
  confs: [default]
    found graphframes#graphframes;0.8.3-spark3.5-s_2.12 in spark-packages
    found org.scala-lang.modules#scala-xml_2.12;1.7.16 in central
```

jupyterhub wordcount-spark (unsaved changes)

File Edit View Insert Cell Kernel Widgets Help Trusted PySpark

```
<pyspark.sql.session.SparkSession object at 0xfffff89dee730>

In [2]: # si esta en EMR o Databricks, estos objetos ya están preconstruidos:
sc
<SparkContext master=yarn appName=livy-session-14>

In [*]: # WORDCOUNT COMPACTO
# en AWS S3
files_rdd = sc.textFile('s3a://emrsamuel/datasets/gutenberg-small/AbrahamLincoln__LincolnLetters.txt')

# en gdrive:
#files_rdd = sc.textFile("gdrive/MyDrive/st0263-252/bigdata/datasets/gutenberg-small/*.txt")

# local:
#files_rdd = sc.textFile("../datasets/gutenberg-small/*.txt")
wc_unsort = files_rdd.flatMap(lambda line: line.split()).map(lambda word: (word, 1)).reduceByKey(lambda a, b: a + b)
wc = wc_unsort.sortBy(lambda a: -a[1])
for tupla in wc.take(10):
    print(tupla)

Progress: [ ]
```

```
In [ ]: # WORDCOUNT PASO A PASO

In [*]: files = sc.textfile("gdrive/MyDrive/st0263-252/bigdata/datasets/gutenberg-small/*.txt")
for f in files.take(10):
    print(f)

In [*]: tokens = files.flatMap(lambda line: line.split())
for t in tokens.take(10):
    print(t)

In [*]: wc1 = tokens.map(lambda word: (word, 1))
for c in wc1.take(10):
    print(c)

In [*]: wc = wc1.reduceByKey(lambda a, b: a + b)
for c in wc.take(10):
    print(c)

In [*]: wcsort = wc.sortBy(lambda a: -a[1])
for c in wcsort.take(10):
    print(c)
```